
CS 152 Laboratory Exercise 2

*Professor: Krste Asanović
TAs: Albert Ou and Yue Dai
Department of Electrical Engineering & Computer Sciences
University of California, Berkeley*

February 21, 2020

Revision History

Revision	Date	Author(s)	Description
1.0	2020-02-21	aou	Initial release

1 Introduction and Goals

The goal of this laboratory assignment is to study **processor memory hierarchy** design by conducting experiments on realistic RISC-V implementations. You will be running simulations of silicon-proven RTL along with a DRAM model.

1.1 Graded Items

All reports are to be submitted through **Gradescope**. Please label each section of the results clearly. All directed items need to be turned in for evaluation. Your group only needs to submit *one* of the problems in the open-ended portion.

- (Directed) Problem 3: Matrix Transposition Case Study
- (Open-ended) Problem 4.1: Validation and Reverse Engineering of Memory Hierarchies
- (Open-ended) Problem 4.2: Design Your Own Hardware Prefetcher
- (Open-ended) Problem 4.3: Design Your Own Replacement Policy and Victim Cache
- (Directed) Problem 5: Feedback

! → Lab reports must be written in *readable* English; avoid raw dumps of logfiles. **Your lab report must be typed, and the open-ended portion must not exceed six (6) pages.** Charts, tables, and figures – where appropriate – are excellent ways to succinctly summarize your data.

2 Background

This lab, as well as subsequent CS 152 labs, is based on the **Chipyard** framework being actively developed UC Berkeley. This enables us to explore more sophisticated hardware designs than the rudimentary Sodor processors from Lab 1.

2.1 Chipyard

Chipyard is an integrated design, simulation, and implementation framework for agile development of systems-on-chip (SoCs). It combines Chisel, the Rocket Chip generator, and other Berkeley projects to produce a full-featured RISC-V SoC from a rich library of processor cores, accelerators, memory system components, and I/O peripherals. Chipyard supports several hardware development flows, including software RTL simulation, FPGA-accelerated simulation (FireSim), and automated VLSI methodologies (Hammer).

! → Chipyard documentation: <https://chipyard.readthedocs.io/en/latest/>

2.2 Rocket Chip

Rocket Chip [1] is an open-source SoC generator originally developed at UC Berkeley. It leverages Chisel to compose a library of highly parameterized generators for cores, caches, and interconnects into an integrated SoC. It has been the basis of numerous silicon-proven designs in both research and industry.

Rocket Chip can generate a practically unbounded space of *instances*, including many parameter sets that are impractical or suboptimal. In this lab, we will examine a variety of design points, each with a different memory hierarchy, to explore the concepts described in class. All Rocket Chip instances used in this lab have three major components: processor cores, a cache hierarchy, and an outer memory system.

2.2.1 Rocket Microarchitecture

Rocket Chip derives its name from the Rocket core that it instantiates by default: a 5-stage, single-issue, in-order RISC-V processor. The instances of Rocket used in this lab implement the RV64IMAFDC instruction set variant¹, which refers to the 64-bit RISC-V base ISA (RV64I) along with a set of useful extensions [2]: M for integer multiply/divide, A for atomic memory operations, F and D for single- and double-precision floating-point, and C for 16-bit compressed representations of common instructions

Rocket also supports the RISC-V privileged architecture [3] with machine, supervisor, and user modes. It has an MMU that implements the Sv39 virtual memory scheme, which provides 39-bit virtual address spaces with 4 KiB pages. As such, these designs are capable of booting mainstream operating systems such as Linux.

Rocket has been extensively optimized for efficient ASIC implementation, resulting in specific microarchitectural adaptations that differentiate it from the classic 5-stage RISC pipeline normally seen in educational settings. In particular, the overall design is mainly

¹ Also known as RV64GC, with G (“general-purpose”) being the canonical shorthand for “IMAFD”




Figure 1: Rocket pipeline

concerned with (1) minimizing high-fanout stall signals and (2) restructuring pipeline logic to cope with long clock-to-Q delays of compiler-generated SRAMs. Several factors contribute to improved reduction in critical paths compared to more naive approaches:

1. Instructions are not permitted to stall except in the ID stage for data and known structural hazards.
2. Most hazards that arise in EX or later stages are handled by replaying (re-fetching and re-executing) the instruction upon reaching WB (not unlike how exceptions propagate down the pipeline). One notable case is load-hit speculation, in which an instruction that depends on a load result can be issued before it is known whether the load is a cache hit.
3. Branch conditions are resolved in EX, but the PC is redirected in MEM. The 3-cycle mispredict penalty is mitigated by branch prediction provided by a configurable branch target buffer (BTB), branch history table (BHT), and a return address stack (RAS).
4. Bypass muxes are moved into EX with the selects precomputed in ID; bypass data comes directly from pipeline registers to the extent possible.
5. Some variable-latency operations (e.g., L1D miss, divide) use a scoreboard to track pending register writes. This enables instructions to complete out of program order so that a long-latency operation does not halt the pipeline for subsequent instructions. Consequently, with a non-blocking L1 data cache, multiple misses can be serviced simultaneously.

2.2.2 Cache Hierarchy

The basic unit of replication for a core in Rocket Chip is a *tile*.² Each tile consists of one core (Rocket) and a portion of the inner cache hierarchy that is private to each core:

- L1 instruction cache (L1I)
- L1 data cache (L1D) of either a blocking or a non-blocking design
- fully-associative L1 instruction and data TLBs

² Although Rocket Chip can generate multi-core instances, this lab will feature only single-tile instances.

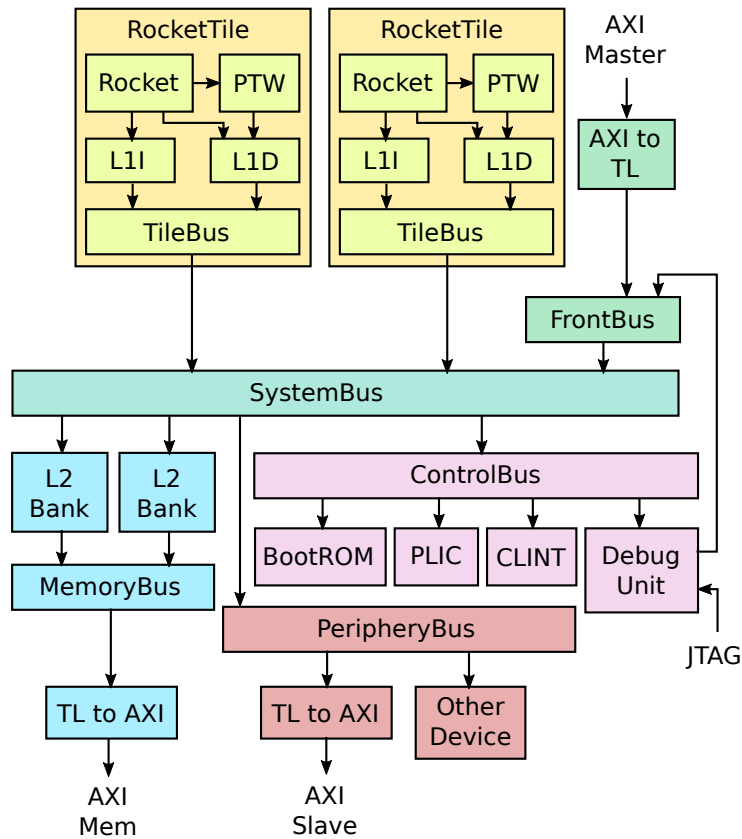


Figure 2: A generic Rocket Chip instance

- optional unified direct-mapped L2 TLB
- hardware page table walker

SoC instances can optionally be configured with a unified, inclusive, multi-banked L2 cache as a last-level cache shared between tiles. If an L2 cache is not present, an L2 broadcast hub is instantiated in its place to maintain coherence between the L1 caches. Each of these structures exposes various parameters such as capacity, associativity, replacement policy, and cache line size, which are set through a Scala-based configuration system at elaboration time.

2.2.3 Outer Memory System

The L2 coherence agent (either the L2 cache or broadcast hub) makes requests to an outer memory system through a AXI4 master port. This top-level port would typically interface with a DRAM controller, but since an actual DRAM controller implementation is not openly available, we instead attach a model that simulates the functional and timing behaviors of a DDR3 memory system. The default SoC configuration presents a single memory channel, but the system can be configured to use multiple channels for greater bandwidth.

2.3 FASED: Modeling Memory Systems In FireSim

FireSim [4] is an open-source cycle-accurate FPGA-accelerated full-system hardware simulation platform that runs on cloud FPGAs, such Amazon EC2 F1. Although we will not be using FPGAs for this lab, a large portion of its infrastructure will be repurposed for *software* RTL simulation, notably FASED [5], a parameterized generator of composable, high-fidelity, synthesizable last-level-cache and DRAM models.

3 Directed Portion (30%)

3.1 Terminology and Conventions

Throughout this course, the term *host* refers to the machine on which the simulation runs, while *target* refers to the machine being simulated. For this lab, an instructional server will act as the host, and the RISC-V processors will be the target machines.

UNIX shell commands to be run on the host are prefixed with the prompt “eecs\$”.

3.2 Setup

To complete this lab, `ssh` into an instructional server with the instructional computing account provided to you. The lab infrastructure has been set up to run on the `eda{1..8}.eecs.berkeley.edu` machines (`eda-1.eecs`, `eda-2.eecs`, etc.).

Once logged in, source the following script to initialize your shell environment so as to be able to access to the tools for this lab. Run it before each session.

```
eecs$ source ~cs152/sp20/cs152.lab2.bashrc
```

First, clone the lab materials into an appropriate workspace and initialize the submodules.

- ! → **It is highly recommended** to work in the local `/scratch` partition to avoid issues with filesystem performance and quotas. Even simulations of modest length (few hundred thousand cycles) can produce a few gigabytes of logs and waveform dumps. Do not use your NFS home directory to avoid slowing down the simulation. Remember that `/scratch` is *not* backed up automatically.

```
eecs$ mkdir -p 0700 /scratch/${USER}
eecs$ cd /scratch/${USER}
eecs$ git clone ~cs152/sp20/lab2.git
eecs$ LAB2ROOT="$(pwd)"
eecs$ ./scripts/init-submodules-no-riscv-tools.sh
```

The remainder of this lab will use `${LAB2ROOT}` to denote the path of the `lab2` working tree. Its directory structure is outlined below:

`${LAB2ROOT}`

lab/

directed/

[Source code for Lab 2 directed portion](#)

open1/	Source code and tools for Problem 4.1
open2/	Source code and tools for Problem 4.2
open3/	Source code and tools for Problem 4.3
generators/	Library of RTL generators
cs152/	SoC configurations for CS152 Lab 2
rocket-chip/	Rocket Chip generator
sifive-blocks/	Miscellaneous peripherals from SiFive
sifive-cache/	Open-source inclusive L2 cache from SiFive
testchipip/	RTL blocks for interfacing with test chips
...	
sims/	
firesim/	FireSim simulation platform
...	
tools/	
chisel3/	Chisel hardware description library
firrtl/	RTL intermediate representation library
barstools/	Collection of common FIRRTL transformations
...	

3.3 Matrix Transposition Case Study

The directed portion will lead you through a simple case study of a matrix transposition kernel with these objectives:

- Illustrate some basic cache optimization techniques
- Conduct a brief design-space exploration of cache configurations using the Rocket Chip parameterization system
- Familiarize you with the RTL simulation flow

We begin with a naive implementation of matrix transposition in `${LAB2ROOT}/lab/directed/transpose.c` that is derived directly from the mathematical definition. Take a moment to understand the source code. Note that both 128×128 input and output matrices are stored in row-major order.

Compile it into a bare-metal binary:

```
eecs$ cd ${LAB2ROOT}/lab/directed
eecs$ make
```

Next, navigate to the FireSim simulation directory and build the simulator. Notice that the `TARGET_CONFIG` variable selects the top-level SoC design to generate – what exactly

this means will be described in the next section. The Chisel design is elaborated into Verilog RTL, which is then compiled into a cycle-accurate simulator.

! → From here on, we will use `${SIMDIR}` to denote the base directory of the FireSim simulation environment.

```
eecs$ SIMDIR=${LAB2ROOT}/sims/firesim/sim
eecs$ cd ${SIMDIR}
eecs$ make TARGET_CONFIG=CS152RocketChipConfig vcs
```

This particular configuration contains a 4 KiB direct-mapped L1 data cache and a 4 KiB direct-mapped L1 instruction cache, both with 64-byte cache lines.

Next, run the naive matrix transposition kernel using the predefined Makefile recipe to invoke the simulator:

```
eecs$ make TARGET_CONFIG=CS152RocketChipConfig run-transpose
```

The program also prints a snapshot of several hardware performance counters in the processor.³ Use this information along with the simulation trace⁴ in `${SIMDIR}/output/f1/*-CS152RocketChipConfig-*/transpose.out` to answer the following questions:

- (3.3.a) How many cycles does the transpose operation take?
- (3.3.b) What is the hit rate of the L1 data cache?
- (3.3.c) Why does the naive transpose code result in non-ideal cache performance?
- (3.3.d) Which memory access in the code incurs the most misses, and why?

3.4 Cache Blocking

Rewrite the transpose code to employ *cache blocking* (loop tiling) using $B \times B$ blocks. Experiment with a few values of B to determine which factor yields the best performance.

- (3.4.a) What is the performance improvement using cache blocking?
- (3.4.b) What is the largest B that one would expect to fit in the 4 KiB L1 cache? Does this result in the highest performance? (If not, why not?)

3.5 Cache Parameters

Navigate to `${LAB2ROOT}/generators/cs152/src/main/scala/TargetConfigs.scala` and examine the definition of `CS152RocketChipConfig`.

In Rocket Chip, a *Config* is a Scala class that sets one or more generator parameters to specific values. Configs are additive, can override each other, and can be composed of

³ Note that in-flight and recently retired instructions may or may not be reflected when reading the performance counters.

⁴ These prints show signals from Rocket's writeback stage; refer to line 898 of `${LAB2ROOT}/generators/rocket-chip/src/main/scala/rocket/RocketCore.scala` to identify each field.

other Configs. `CS152RocketChipConfig` is an example of a Config that combines other Configs through the `++` operator. The constituent Configs are applied from right to left (or from bottom to top) in the chain, by reverse order of precedence. Thus, a Config appearing to the left of (or above) another Config overrides any parameters previously set by the latter. For more information on the Rocket Chip parameter system, read through the Chipyard documentation.⁵

In `CS152RocketChipConfig`, change the associativity of the L1 data cache to 2 by modifying `WithL1DWays` parameter, while also adjusting `WithL1DSets` to keep the overall cache size constant.

Rebuild the simulator and re-run your blocked matrix transpose version. Repeat this with 4 and 8 ways.

(3.5.a) How does performance and miss rate change when associativity is increased?

(3.5.b) Explain why higher associativity is (or is not) beneficial for this particular kernel.

3.6 Multi-level Caches

We will continue our experimentation with the `CS152RocketChipL2Config` design, which is derived from `CS152RocketChipL2Config` but adds a 64 KiB 8-way L2 cache. (Remember to change the L1 configuration back to a 4 KiB direct-mapped cache.)

Modify your transposition code to introduce another level of cache blocking for the L2. Simulate on `CS152RocketChipL2Config` and answer the following:

(3.6.a) What is the optimal blocking factor B' for the L2?

(3.6.b) What is the performance improvement with this change?

Change the `WithCacheBlockBytes` parameter to have 128-byte cache lines instead of 64-byte cache lines. (Adjust the L1 parameters accordingly to keep the cache size constant.)

(3.6.c) How does the cache line size affect performance and the choice of the blocking factors B and B' ? Why?

⁵<https://chipyard.readthedocs.io/en/latest/Chipyard-Basics/Configs-Parameters-Mixins.html>

4 Open-ended Portion (70%)

Select *one* of the following questions per team. The open-ended portion is worth a considerable fraction of the grade for the lab, and the grade depends on how comprehensive the process to your conclusion is.

4.1 Validation and Reverse Engineering of Memory Hierarchies

In this problem, we will try to infer fundamental parameters of a memory system by running user code and measuring execution latency. This is useful for a variety of reasons:

- To help guide application optimizations when the underlying microarchitecture is unknown or undisclosed.
- To validate memory system performance before tape-out. Some of the most insidious bugs in computer system design are performance bugs, since applications still execute correctly but only more slowly. We would like to catch these bugs before committing a design to silicon, but without a performance model of the target machine, they may go undiscovered.
- Using the same principle as above, to help validate simulation models. FASED, like many simulation models, splits its timing and functional modeling. This approach makes it possible to build highly complex cache and memory models – but it is very easy to write “correct” yet fundamentally broken timing models.⁶

! → We have provided a mystery Rocket Chip configuration for you to characterize, aptly named `CS152RocketChipMysteryConfig`.

4.1.1 Cache Sizes and Access Latency

We will first run the `caches` micro-benchmark, which comes from the `ccbench` suite developed by Christopher Celio, to determine the cache sizes and access latency at each level of the cache hierarchy.

`caches` executes a single-threaded pointer chase on an array of a given size, in which each 4-byte array element yields the index of the next element to access.

```
start_cycles = get_cycles();
for (uint32_t k = 0; k < g_num_iterations; k++) {
    idx = arr_n_ptr[idx];
}
end_cycles = get_cycles();
```

As the array exceeds the size of a cache, the sharp increase in misses is observable from the longer time that the benchmark takes to run. To accurately isolate the load latency of an individual element, both spatial locality and memory prefetching must be defeated. This involves striding the indices to point to different cache lines and randomly sorting

⁶ For example, in one FASED bug, a write address was mistakenly being used for a read access to the tag array of the last-level cache model. While this would cause a real cache to return incorrect data, it simply manifested as small timing aberration in the model.

the indices within a virtual page (so TLB locality is still maintained).

We have provided a Makefile to automate running ccbench in simulation and visualizing the result. First, sweep across a predefined range of array sizes:

```
eecs$ cd ${LAB2ROOT}/lab/open1
eecs$ make ccbench-sweep
```

This should require approximately **an hour** to finish. Each benchmark run outputs a log statement in the form:

```
App:[caches],NumThreads:[0],AppSize:[1024],Time:[4.01151],TimeUnits:[
↪ Cycles Per Iteration],NumIterations:[100000]
```

“AppSize” records the array size in terms of 4-byte elements, and “Time” records the average cycles spent per iteration. Run the following recipe to extract these lines into a consolidated report file and invoke the plotting script from ccbench:

```
eecs$ make ccbench-plot
```

This generates a plot of cycles per iteration versus array size. Open `${LAB2ROOT}/lab/open1/ccbench/caches/plots/plot-CS152RocketChipMysteryConfig.pdf` and answer the following:

- L1 D cache size
- L1 D cache latency
- L2 cache size
- L2 cache latency
- DRAM latency

Save the plot for your report.

4.1.2 Other Parameters

Finally, try to empirically deduce some more subtle parameters:

- L1 D cache associativity (or equivalently, cache line size)
- L1 I cache size
- L1 I cache associativity
- L1 I cache replacement policy
- L1 I TLB reach
- L1 D TLB reach
- L2 TLB reach
- L2 TLB hit latency
- DRAM page policy (open or closed)
- Aggregate DRAM page size (ranks \times banks)
- Number of DRAM ranks (you can assume there are 8 banks)
- DRAM CAS/RCD/RP latencies (you can safely assume they are identical)

You may either modify the `caches` benchmark, for which the source code can be found in `${LAB2ROOT}/lab/open1/ccbench/caches/caches.c`, or write your own program.

For example, to manually compile a simple C user program and run it in simulation:

```
eecs$ SIMDIR=${LAB2ROOT}/sims/firesim/sim
eecs$ riscv64-unknown-elf-gcc -O2 -o ${SIMDIR}/example -o example.c
eecs$ cd ${SIMDIR}
eecs$ make TARGET_CONFIG=CS152RocketChipMysteryConfig BINARY="./example" run-pk
```

Keep in mind that your program executes in user mode with virtual addressing enabled, so generally you cannot access arbitrary memory locations that have not been properly allocated, either statically or dynamically with `malloc()` or `mmap()`.

For a rudimentary timer, the `cycle` CSR (which counts the number of cycles after reset) can be read with inline assembly:

```
static inline unsigned long get_cycles(void) {
    unsigned long cycles;
    __asm__ __volatile__ ("rdcycle %0" : "=r" (cycles));
    return cycles;
}
```

Feel free to test your code on other known configurations, such as the one used in the directed portion or any new ones that you define.

4.1.3 Submission

For the given `CS152RocketChipMysteryConfig` instance, report the cache capacities, access latencies, and block sizes as indicated by the `caches` benchmark. (Include the plots from `ccbench`.)

Then provide your best estimate for as many of the other parameters as possible – aim for *at least six*. For each of those parameters, explain how you measured it, referring to your code as necessary. If you are not certain that you can accurately determine enough parameters, still provide your code and explain what you tried. More credit will be awarded for a measured and analyzed negative result than an ill-justified guess which might be correct by coincidence. If you have data or plots to show that your code works on known instances, include them in your justification.

Feel free reach out to your GSI if you need help understanding `ccbench`, Rocket Chip, or anything else regarding this problem.

4.2 Design Your Own Hardware Prefetcher

In this problem, we will build a hardware prefetcher (in either Chisel or C++) in hopes of improving the performance on various benchmarks.

We will use a Rocket Chip instance that has a 16 KiB 4-way set-associative L1 data cache and 64-byte cache lines. Unlike the directed portion, this configuration uses the *non-blocking data cache*, so that Rocket can take advantage of hit-under-miss while a prefetch is being serviced.




Figure 3: L1 prefetcher integration

4.2.1 Interfaces

Partly for convenience and modularity, the prefetcher is integrated with the tile as a separate module from the cache itself. It can be considered another client of the data cache, like the core.

Start by navigating to `${LAB2ROOT}/generators/rocket-chip/src/main/scala/rocket/L1Prefetcher.scala`. This Chisel file contains our generator framework for creating prefetchers. All prefetcher modules inherit from the `L1Prefetcher` base class, which specifies a common set of I/O ports (grouped together in the `L1PrefetcherIO` bundle).

To observe the stream of memory requests and misses, the prefetcher snoops on the `HellaCacheIO` interface between the core and L1 data cache. A curated subset of signals from the `HellaCacheReq` bundle is presented to the prefetcher:

Input Signal	Description
<code>io.cpu.req.valid</code>	Asserted when Rocket's execute stage sends a request to the L1D cache
<code>io.cpu.req.bits.addr</code>	Virtual address of the access
<code>io.cpu.req.bits.cmd</code>	Memory operation type (e.g., 0=load, 1=store, etc.)
<code>io.cpu.req.bits.size</code>	Logarithm of access size (e.g., 0=1 byte, 3=8 bytes)
<code>io.cpu.miss</code>	Asserted when a cache miss is being reported to Rocket's writeback stage (Note: This is delayed by two cycles after the original request)

The prefetcher has a simplified outgoing interface through which it can inject prefetch requests into the L1D. This is a decoupled interface that uses ready/valid hand-shaking⁷ to coordinate the source and sink: Both `io.dmem.valid` and `io.dmem.ready` must be high during the same cycle to initiate a prefetch.

⁷ <https://inst.eecs.berkeley.edu/~cs150/Documents/Interfaces.pdf>

Input Signal	Description
<code>io.dmem.ready</code>	Asserted when the L1D can accept a request
Output Signal	Description
<code>io.dmem.valid</code>	Indicates that the request is valid
<code>io.dmem.bits.addr</code>	Virtual address to prefetch (Note: This must be aligned to an 8-byte boundary)
<code>io.dmem.bits.store</code>	Indicates intent to write

The single port to the L1D is arbitrated between the core, prefetcher, and the page table walker, with the prefetcher being given the lowest priority so as to avoid blocking actual memory requests. However, for more aggressive prefetching schemes, it may be desirable to implement some form of throttling to ensure that prefetches do not excessively occupy the MSHRs (miss status handling registers). For example, prefetches could be rate-limited based on a fixed interval or a feedback loop that adapts to miss rate.

For reference, the `ExampleL1Prefetcher` module is provided as a demonstration on how to use the `L1PrefetcherIO` interface described above. This is a naive implementation of the one-block prefetch-on-miss scheme from lecture, but its simplistic design actually turns out to be quite ineffectual, causing a moderate performance degradation more often than it helps. Hopefully yours is a superior solution!

- ! → Once you understand how the interfaces work, implement your own hardware prefetcher within the empty `CustomL1Prefetcher` module provided in `L1Prefetcher.scala` (search for a TODO comment).

4.2.2 C++ Modeling

As an alternative to writing Chisel, you also have the option of implementing your prefetcher in C++ as a software model that is co-simulated with Rocket Chip. This method uses SystemVerilog DPI (Direct Programming Interface) to enable a Verilog wrapper module to call C++ functions.

Start by navigating to `${LAB2ROOT}/generators/rocket-chip/src/main/resources/csrc/L1Prefetcher.cc`. This C++ file contains two functions to work with:

Function	Description
<code>prefetcher_init()</code>	This is called at the beginning of simulation and can be used to initialize global state.
<code>prefetcher_tick()</code>	This is called for each target clock cycle and is where the bulk of your prefetcher logic will reside. The function signature matches the I/O ports described above. Values for output signals (<code>dmem_*</code>) are assigned by dereferencing the argument pointers.

4.2.3 Building

To first use your custom hardware prefetcher, you must modify the `WithL1Prefetcher` Config defined in `${LAB2ROOT}/generators/cs152/src/main/scala/TargetConfigs.scala`. Replace the default `ExampleL1Prefetcher` instantiation with the appropriate module, either `CustomL1Prefetcher` (Chisel) or `ModelL1Prefetcher` (C++), like so:

```
class WithL1Prefetcher extends Config((site, here, up) => {
  case BuildL1Prefetcher =>
    Some((p: Parameters) => Module(new CustomL1Prefetcher()(p)))
})
```

A specific top-level configuration for this problem (`CS152RocketChipPrefetchConfig`) has already been prepared for you, which includes the `WithL1Prefetcher` Config. To build the simulator:

```
eeecs$ SIMDIR=${LAB2ROOT}/sims/firesim/sim
eeecs$ cd ${SIMDIR}
eeecs$ make TARGET_CONFIG=CS152RocketChipPrefetchConfig vcs
```

4.2.4 Simulating

First see how your prefetcher performs on the matrix transposition kernel from the directed portion. You may choose to use the naive code or your L1 cache-blocked version.

```
eeecs$ cd ${SIMDIR}
eeecs$ make TARGET_CONFIG=CS152RocketChipPrefetchConfig run-transpose
```

Next test your prefetcher on the suite of benchmarks from Lab 1:

```
eeecs$ make TARGET_CONFIG=CS152RocketChipPrefetchConfig run-bmark-tests
```

Lastly, as an example of a more complex application, we have also included the Graph Algorithm Performance Benchmark Suite (GAPBS) [6], which consists of portable, high-performance implementations for six fundamental graph algorithms developed by Scott Beamer. Specifically, our focus is on the direction-optimizing variant of Breadth-First Search (BFS). Smaller inputs (Kronecker graphs with 2^{10} vertices) will be used here, as the reference inputs such as real social network graphs are too intensive in memory requirements and simulation time.

```
eeecs$ cd ${LAB2ROOT}/lab/open2
eeecs$ make
eeecs$ cd ${SIMDIR}
eeecs$ make TARGET_CONFIG=CS152RocketChipPrefetchConfig run-bfs
```

You can find the simulation traces in `${SIMDIR}/output/f1/<tuple>/*.out`, where `<tuple>` is `FireSimCS152-CS152RocketChipPrefetchConfig-BaseF1Config`.

4.2.5 Debugging

To dump waveforms from simulation, run the debug versions of the `make` targets:

```
eecs$ make TARGET_CONFIG=CS152RocketChipPrefetchConfig vcs-debug
eecs$ make TARGET_CONFIG=CS152RocketChipPrefetchConfig run-transpose-debug
eecs$ make TARGET_CONFIG=CS152RocketChipPrefetchConfig run-bmark-tests-debug
eecs$ make TARGET_CONFIG=CS152RocketChipPrefetchConfig run-bfs-debug
```

The waveforms are written to `${SIMDIR}/output/f1/<tuple>/*.vpd`. The prefetcher instance is found under `emul.F1Shim.top.sim.target.FireSimCS152.targets.tile.prefetchOpt` in the module hierarchy. Waveforms can be viewed on the instructional servers with the DVE application (requires X11 forwarding over `ssh` or `X2Go`):

```
eecs$ dve & # '&' backgrounds the process
```

One peculiarity of FireSim is that these RTL simulations are technically simulating a simulator hosted on an FPGA⁸, a necessary sacrifice to be able to utilize the DRAM timing model. Consequently, the simulation of target design does not always advance at every clock cycle, which manifests in the waveform as an arbitrarily gated clock signal.

4.2.6 Submission

Report performance metrics and cache statistics from running the various benchmarks with the prefetcher enabled. Compare these to results gathered from running on the baseline `CS152RocketChipNoPrefetchConfig` system, which omits the prefetcher but is otherwise identical. Include the source code for your implementation in an appendix.

In your report, describe your design and any implementation challenges in detail. Here are some suggestions to consider in your evaluating:

- What memory access patterns or instances of locality were you attempting to capture? Were you tuning for any specific program behaviors?
- Explain your design rationale and the various approaches that you considered. What worked and what did not?
- Analyze the impact on miss rate and CPI. Were any results surprising?
- Estimate the cost in hardware resources or cycle time for your implementation.
- What changes or alternative approaches would you pursue as future work if more time were available?

A negative result is perfectly acceptable so long as you reason about why the outcomes differed from your expectations. (Designing an effective prefetcher is a non-trivial task! This exercise is partly meant to underscore the challenges of prototyping an idea.)

Feel free reach out to your GSI if you need help understanding Chisel, Rocket Chip, or

⁸<http://docs.fires.im/en/latest/Advanced-Usage/Debugging/RTL-Simulation.html>

anything else regarding this problem.

4.3 Design Your Own Replacement Policy and Victim Cache

For this problem, we would like to investigate whether a different cache replacement policy, combined with a victim cache, would improve the performance of five selected SPEC benchmarks compared to random replacement.

Assume you are designing for a 16 KiB 4-way set-associative L1 data cache, where the backside is connected to DRAM. You will model your cache modifications in **spike**, a functional ISA simulator for RISC-V that has been extended with a basic cache model. The simulator feeds memory addresses through a simulated cache (with a given size, associativity, and block size) to compute the number of accesses, hits, and misses. While **spike** does *not* model microarchitectural timings and is therefore not cycle-accurate, its speed lets us execute much longer programs ordinarily infeasible in RTL simulation.

- ! → The only constraints are that you can only add *less than 2^{11} bits* of state (in either flip-flops or SRAM) to support your new replacement policy and *less than 2^{13} bits* for the victim cache. Assume that physical addresses are 56 bits wide.

4.3.1 SPEC CPU2006

The SPEC CPU2006 package is a former⁹ industry-standard benchmark suite for evaluating general-purpose processors, memory systems, and compilers [7]. You will be running five benchmarks from SPECint (integer) and SPECfp (floating-point) on smaller *test* inputs.¹⁰ Brief descriptions of them, taken from the SPEC documentation, follow:

- 401.bzip2 is based on bzip2 version 1.0.3, modified to perform compression and decompression in memory instead of file I/O.
- 429.mcf is derived from MCF, a program used for single-depot vehicle scheduling in public mass transportation. It features a specialized version of the simplex algorithm for network flow problems.
- 450.soplex is based on SoPlex 1.2.1. It solves a linear program using a simplex algorithm and sparse linear algebra.
- 458.sjeng is based on Sjeng 11.2, a program that plays chess and several chess variants. It attempts to find the best move via a combination of alpha-beta or priority proof number tree searches, advanced move ordering, positional evaluation and heuristic forward pruning.
- 470.lbm implements the “Lattice Boltzmann Method” to simulate incompressible fluids in 3D.

⁹ It has since been replaced by SPEC CPU2017; however, we still opt to use SPEC CPU2006 for being much simpler to cross-compile for RISC-V.

¹⁰ Each benchmark with reference inputs generally require a day to run on a typical Rocket Chip instance mapped to an FPGA. In this case, we are more interested in stressing the caches than reporting valid benchmark scores, so minor adjustments have been made to limit simulation time.

To build **spike** and simulate all benchmarks using its cache model:

```
eeecs$ cd ${LAB2ROOT}/lab/open3
eeecs$ make run
```

The benchmarks should take around six minutes to execute overall. For quicker testing, individual programs can be re-run with **make run-*X***, where *X* is the name of the benchmark without the numerical prefix (e.g., **run-bzip2**, **run-mcf**, etc.).

The simulation output is recorded in `${LAB2ROOT}/lab/open3/CPU2006/CPU2006/build.riscv/*.out`, and the compiled SPEC binaries can also be found in that same directory. Removing these `*.out` files forces **make** to re-run the simulations later:

```
eeecs$ make clean-run
```

And to purge all generated files to trigger a full rebuild:

```
eeecs$ make clean
```

4.3.2 Modifying the Cache Simulator

Navigate to `${LAB2ROOT}/lab/open3/riscv-isa-sim/riscv/cachesim.cc`, where you will find the definition for the `cache_sim_t` C++ class instantiated in **spike**. After taking some time to understand how the current cache model operates, modify the `cache_sim_t::victimize()` function to implement your custom replacement policy.

To rebuild the simulator without running any benchmarks:

```
eeecs$ cd ${LAB2ROOT}/lab/open3
eeecs$ make spike
```

4.3.3 Adding a Victim Cache

Generally, the associativity of a cache (number of ways) presents a trade-off between access time and conflict misses. In order to reduce conflict misses without affecting access times, N. Jouppi proposed *victim caching* [8] in which a small fully-associative secondary cache, called a victim cache, is added to a direct-mapped L1 cache to hold recently evicted cache lines.

We are interested in whether a victim cache would be worthwhile to implement even for a set-associative cache with a moderate number of ways. To get a sense of the scope for potential improvement, you may want to first augment `cache_sim_t` to track the number of conflict misses. Then implement your own victim caching scheme inside the cache model. You will likely need to modify the `cache_sim_t::access()` function, which calls `cache_sim_t::victimize()`.

4.3.4 Submission

In your report, describe how your cache replacement policy and victim cache work, using visual aids (e.g., block diagrams) where appropriate to illustrate their operation. Include a diff of your modifications to the cache model in an appendix.

Report cache statistics from running the SPEC benchmark suite on your modified cache, and compare them to the original cache. Try to explain your results as best you can. Here are some suggestions to consider in your analysis on the effectiveness of your design:

- Based on the number of memory accesses, misses, and instructions retired, what effect on **AMAT** and **CPI** do you think your new cache would have?
- Estimate the **cost in resources or cycle time** if you were to implement your design in hardware. How much **additional state would be required in terms of bits?**
- For the given set of benchmarks, how problematic are **conflict misses compared to compulsory or capacity misses?** Would the addition of a victim cache be **justifiable?**
- What is the **minimum number of victim cache entries** that you would recommend? Are there **diminishing returns to increasing victim cache size?**
- Program behavior can sometimes be characterized by distinct execution “phases”. Does the **miss rate vary over time** within a benchmark?

Feel free to reach out to your GSI if you need help understanding the ISA simulator, the cache model, or anything else regarding this problem.

5 Feedback Portion

In order to improve the labs for the next offering of this course, we would like your feedback. Please append your feedback to your individual report for the directed portion.

- How many hours did the directed portion take you?
- How many hours did you spend on the open-ended portion?
- Was this lab boring?
- What did you learn?
- Is there anything that you would change?

Feel free to write as much or as little as you prefer (a point will be deducted only if left completely empty).

5.1 Team Feedback

In addition to feedback on the lab itself, please answer a few questions about your team:

- In a few sentences, describe your contributions to the project.
- Describe the contribution of each of your team members.
- Do you think that every member of the team contributed fairly? If not, why?

6 Acknowledgments

This lab was heavily inspired by the previous set of CS 152 labs developed by Henry Cook, Yunsup Lee, and Andrew Waterman, which targeted functional simulators such as Simics and Spike. More recent iterations of this lab, developed by Donggyu Kim, David Biancolin, and Albert Magyar, used FireSim to run FPGA-based simulations on Amazon EC2 F1.

References

- [1] K. Asanović, R. Aviúzienis, J. Bachrach, S. Beamer, D. Biancolin, C. Celio, H. Cook, D. Dabbelt, J. Hauser, A. Izraelevitz, S. Karandikar, B. Keller, D. Kim, J. Koenig, Y. Lee, E. Love, M. Maas, A. Magyar, H. Mao, M. Moreto, A. Ou, D. A. Patterson, B. Richards, C. Schmidt, S. Twigg, H. Vo, and A. Waterman, “The Rocket Chip generator,” EECS Department, University of California, Berkeley, Tech. Rep. UCB/EECS-2016-17, Apr. 2016.
- [2] A. Waterman and K. Asanović, Eds., *The RISC-V instruction set manual, volume I: User-level ISA*, Version 20191213, RISC-V Foundation, Dec. 2019. [Online]. Available: <https://riscv.org/specifications/>.
- [3] A. Waterman and K. Asanović, Eds., *The RISC-V instruction set manual, volume II: Privileged architecture*, Version 20190608-Priv-MSU-Ratified, RISC-V Foundation, Jun. 2019. [Online]. Available: <https://riscv.org/specifications/privileged-isa/>.
- [4] S. Karandikar, H. Mao, D. Kim, D. Biancolin, A. Amid, D. Lee, N. Pemberton, E. Amaro, C. Schmidt, A. Chopra, Q. Huang, K. Kovacs, B. Nikolić, R. Katz, J. Bachrach, and K. Asanović, “FireSim: FPGA-accelerated cycle-exact scale-out system simulation in the public cloud,” in *Proceedings of the 45th Annual International Symposium on Computer Architecture*, Los Angeles, California: IEEE Press, Jun. 2018, 2942. DOI: 10.1109/ISCA.2018.00014.
- [5] D. Biancolin, S. Karandikar, D. Kim, J. Koenig, A. Waterman, J. Bachrach, and K. Asanović, “FASED: FPGA-accelerated simulation and evaluation of DRAM,” in *Proceedings of the 2019 ACM/SIGDA International Symposium on Field-Programmable Gate Arrays*, Seaside, CA, USA: Association for Computing Machinery, Feb. 2019, 330339. DOI: 10.1145/3289602.3293894.
- [6] S. Beamer, K. Asanović, and D. Patterson, *The GAP benchmark suite*, 2015. arXiv: 1508.03619 [cs.DC]. [Online]. Available: <http://gap.cs.berkeley.edu/benchmark.html>.
- [7] J. L. Henning, “SPEC CPU2006 benchmark descriptions,” *SIGARCH Computer Architecture News*, vol. 34, no. 4, 117, Sep. 2006. DOI: 10.1145/1186736.1186737.
- [8] N. P. Jouppi, “Improving direct-mapped cache performance by the addition of a small fully-associative cache and prefetch buffers,” *SIGARCH Computer Architecture News*, vol. 18, no. 2SI, 364373, May 1990. DOI: 10.1145/325096.325162.